

# ML\_with\_TF\_in\_VertexAI

April 25, 2023

## 0.1 Create minimal training and validation data

```
[1]: import os, json, math, shutil
import numpy as np
import tensorflow as tf
#set environment variables to use bash cells
PROJECT = !(gcloud config get-value project)
# print(PROJECT)
PROJECT = PROJECT[0]
REGION = 'us-central1'
BUCKET = "{}-dsongcp".format(PROJECT)
os.environ["ENDPOINT_NAME"] = "flights"
os.environ["BUCKET"] = BUCKET
os.environ["REGION"] = REGION
os.environ["TF_VERSION"] = "2-" + tf.__version__[2:3]

[2]: %%bigquery
# Export files that contain training, validation data
CREATE OR REPLACE TABLE dsongcp.flights_train_data AS
SELECT
  IF(arr_delay < 15, 1.0, 0.0) AS ontime,
  dep_delay,
  taxi_out,
  distance,
  origin,
  dest,
  EXTRACT(hour FROM dep_time) AS dep_hour,
  IF (EXTRACT(dayofweek FROM dep_time) BETWEEN 2 AND 6, 1, 0) AS is_weekday,
  UNIQUE_CARRIER AS carrier,
  dep_airport_lat,
  dep_airport_lon,
  arr_airport_lat,
  arr_airport_lon
FROM dsongcp.flights_tzcorr f
JOIN dsongcp.trainday t
ON f.FL_DATE = t.FL_DATE
WHERE
  f.CANCELLED = False AND
```

```
f.DIVERTED = False AND  
is_train_day = 'True'
```

Query is running: 0%| |

[2]: Empty DataFrame  
Columns: []  
Index: []

```
[3]: %%bigquery  
# Create the evaluation dataset flights_eval_data for model evaluation:  
CREATE OR REPLACE TABLE dsongcp.flights_eval_data AS  
SELECT  
  IF(arr_delay < 15, 1.0, 0.0) AS ontime,  
  dep_delay,  
  taxi_out,  
  distance,  
  origin,  
  dest,  
  EXTRACT(hour FROM dep_time) AS dep_hour,  
  IF (EXTRACT(dayofweek FROM dep_time) BETWEEN 2 AND 6, 1, 0) AS is_weekday,  
  UNIQUE_CARRIER AS carrier,  
  dep_airport_lat,  
  dep_airport_lon,  
  arr_airport_lat,  
  arr_airport_lon  
FROM dsongcp.flights_tzcorr f  
JOIN dsongcp.trainday t  
ON f.FL_DATE = t.FL_DATE  
WHERE  
  f.CANCELLED = False AND  
  f.DIVERTED = False AND  
  is_train_day = 'False'
```

Query is running: 0%| |

[3]: Empty DataFrame  
Columns: []  
Index: []

```
[4]: %%bigquery  
# Create the full dataset flights_all_data using the following code:  
CREATE OR REPLACE TABLE dsongcp.flights_all_data AS  
SELECT  
  IF(arr_delay < 15, 1.0, 0.0) AS ontime,  
  dep_delay,
```

```

taxi_out,
distance,
origin,
dest,
EXTRACT(hour FROM dep_time) AS dep_hour,
IF (EXTRACT(dayofweek FROM dep_time) BETWEEN 2 AND 6, 1, 0) AS is_weekday,
UNIQUE_CARRIER AS carrier,
dep_airport_lat,
dep_airport_lon,
arr_airport_lat,
arr_airport_lon,
IF (is_train_day = 'True',
    IF (ABS(MOD(FARM_FINGERPRINT(CAST(f.FL_DATE AS STRING))), 100)) < 60,
    ↪ 'TRAIN', 'VALIDATE'),
    'TEST') AS data_split
FROM dsongcp.flights_tzcorr f
JOIN dsongcp.trainday t
ON f.FL_DATE = t.FL_DATE
WHERE
    f.CANCELLED = False AND
    f.DIVERTED = False

```

Query is running: 0%| |

[4]: Empty DataFrame  
Columns: []  
Index: []

```

[5]: %%bash
PROJECT=$(gcloud config get-value project)
for dataset in "train" "eval" "all"; do
    TABLE=dsongcp.flights_${dataset}_data
    CSV=gs://${BUCKET}/ch9/data/${dataset}.csv
    echo "Exporting ${TABLE} to ${CSV} and deleting table"
    bq --project_id=${PROJECT} extract --destination_format=CSV $TABLE $CSV
    bq --project_id=${PROJECT} rm -f $TABLE
done

```

Exporting dsongcp.flights\_train\_data to gs://qwiklabs-gcp-01-c3764cc81f6b-dsongcp/ch9/data/train.csv and deleting table  
Exporting dsongcp.flights\_eval\_data to gs://qwiklabs-gcp-01-c3764cc81f6b-dsongcp/ch9/data/eval.csv and deleting table  
Exporting dsongcp.flights\_all\_data to gs://qwiklabs-gcp-01-c3764cc81f6b-dsongcp/ch9/data/all.csv and deleting table

Waiting on bqjob\_r60ae91da6b8e0a89\_00000187b9b16e17\_1 ... (65s) Current status: DONE

Waiting on bqjob\_r7f616cecbe84f130\_00000187b9b27de1\_1 ... (23s) Current status:

DONE

Waiting on bqjob\_r3b286050cf254292\_00000187b9b2e85b\_1 ... (85s) Current status:

DONE

```
[6]: !gsutil ls -lh gs://{BUCKET}/ch9/data
```

```
445.01 MiB  2023-04-25T18:36:38Z  gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/data/all.csv  
115.19 MiB  2023-04-25T18:35:15Z  gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/data/eval.csv  
296.96 MiB  2023-04-25T18:34:37Z  gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/data/train.csv  
TOTAL: 3 objects, 898801258 bytes (857.16 MiB)
```

## 0.2 Create the input data

```
[7]: DEVELOP_MODE = True  
NUM_EXAMPLES = 5000*1000
```

```
[8]: training_data_uri = 'gs://{}/ch9/data/train*'.format(BUCKET)  
validation_data_uri = 'gs://{}/ch9/data/eval*'.format(BUCKET)
```

```
[9]: NBUCKETS = 5  
NEMBEDS = 3  
TRAIN_BATCH_SIZE = 64  
DNN_HIDDEN_UNITS = '64,32'
```

```
[10]: if DEVELOP_MODE:  
    train_df = tf.data.experimental.make_csv_dataset(training_data_uri,  
    ↪ batch_size=5)  
    for n, data in enumerate(train_df):  
        numpy_data = {k: v.numpy() for k, v in data.items()}  
        print(n, numpy_data)  
        if n==1: break
```

```
2023-04-25 18:36:53.734706: W  
tensorflow/stream_executor/platform/default/dso_loader.cc:64] Could not load  
dynamic library 'libcuda.so.1'; dLError: libcuda.so.1: cannot open shared object  
file: No such file or directory; LD_LIBRARY_PATH:  
/usr/local/cuda/lib64:/usr/local/ncc12/lib:/usr/local/cuda/extras/CUPTI/lib64  
2023-04-25 18:36:53.734757: W  
tensorflow/stream_executor/cuda/cuda_driver.cc:269] failed call to cuInit:  
UNKNOWN ERROR (303)  
2023-04-25 18:36:53.734785: I  
tensorflow/stream_executor/cuda/cuda_diagnostics.cc:156] kernel driver does not  
appear to be running on this host (tensorflow): /proc/driver/nvidia/version does  
not exist  
2023-04-25 18:36:53.738700: I tensorflow/core/platform/cpu_feature_guard.cc:142]
```

This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: AVX2 FMA

To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.

2023-04-25 18:36:54.857494: I

tensorflow/compiler/mlir/mlir\_graph\_optimization\_pass.cc:185] None of the MLIR Optimization Passes are enabled (registered 2)

```
0 {'ontime': array([1, 1, 1, 1, 1], dtype=int32), 'dep_delay': array([ 3, 15,
-6,  6, -1], dtype=int32), 'taxi_out': array([10, 13, 10, 13, 47], dtype=int32),
'distance': array([ 216,  967,  727, 1020,  337], dtype=int32), 'origin':
array([b'HNL', b'DEN', b'TPA', b'PHX', b'LAX'], dtype=object), 'dest':
array([b'ITO', b'SFO', b'SDF', b'GEG', b'SFO'], dtype=object), 'dep_hour':
array([21,  2, 22, 20, 17], dtype=int32), 'is_weekday': array([1, 0, 1, 1, 1],
dtype=int32), 'carrier': array([b'HA', b'UA', b'WN', b'WN', b'UA'],
dtype=object), 'dep_airport_lat': array([21.317778, 39.861668, 27.975555,
33.434166, 33.9425  ],
      dtype=float32), 'dep_airport_lon': array([-157.92027 , -104.67306 ,
-82.53333 , -112.011665, -118.40806 ],
      dtype=float32), 'arr_airport_lat': array([19.720278, 37.61889 , 38.174168,
47.619167, 37.61889 ],
      dtype=float32), 'arr_airport_lon': array([-155.04834, -122.375  ,
-85.73639, -117.53528, -122.375  ],
      dtype=float32)}
1 {'ontime': array([1, 0, 1, 1, 1], dtype=int32), 'dep_delay': array([18, 29,
-3,  1, -6], dtype=int32), 'taxi_out': array([11, 17, 12, 12, 14], dtype=int32),
'distance': array([ 550, 1020,  284,  500,  224], dtype=int32), 'origin':
array([b'PDX', b'PHX', b'DTW', b'ATL', b'SEA'], dtype=object), 'dest':
array([b'SFO', b'GEG', b'CIU', b'AEX', b'GEG'], dtype=object), 'dep_hour':
array([ 1,  4, 15, 20,  5], dtype=int32), 'is_weekday': array([1, 1, 0, 1, 1],
dtype=int32), 'carrier': array([b'00', b'AA', b'00', b'EV', b'AS'],
dtype=object), 'dep_airport_lat': array([45.58861 , 33.434166, 42.2125  ,
33.636665, 47.45    ],
      dtype=float32), 'dep_airport_lon': array([-122.59695 , -112.011665,
-83.35333 , -84.42778 , -122.31167 ],
      dtype=float32), 'arr_airport_lat': array([37.61889 , 47.619167, 46.25083 ,
31.3275  , 47.619167],
      dtype=float32), 'arr_airport_lon': array([-122.375  , -117.53528 ,
-84.4725  , -92.548615, -117.53528 ],
      dtype=float32)}
```

```
[11]: def features_and_labels(features):
      label = features.pop('ontime')
      return features, label

      def read_dataset(pattern, batch_size, mode=tf.estimator.ModeKeys.TRAIN,
      ↪truncate=None):
```

```

dataset = tf.data.experimental.make_csv_dataset(pattern, batch_size,
↳num_epochs=1)
dataset = dataset.map(features_and_labels)
if mode == tf.estimator.ModeKeys.TRAIN:
    dataset = dataset.shuffle(batch_size*10)
    dataset = dataset.repeat()
dataset = dataset.prefetch(1)
if truncate is not None:
    dataset = dataset.take(truncate)
return dataset

if DEVELOP_MODE:
    print("Checking input pipeline")
    one_item = read_dataset(training_data_uri, batch_size=2, truncate=1)
    print(list(one_item)) # should print one batch of 2 items

```

Checking input pipeline

```

[(OrderedDict([('dep_delay', <tf.Tensor: shape=(2,), dtype=int32,
numpy=array([54, 1], dtype=int32)>), ('taxi_out', <tf.Tensor: shape=(2,),
dtype=int32, numpy=array([11, 22], dtype=int32)>), ('distance', <tf.Tensor:
shape=(2,), dtype=int32, numpy=array([1846, 262], dtype=int32)>), ('origin',
<tf.Tensor: shape=(2,), dtype=string, numpy=array([b'ORD', b'SBA'],
dtype=object)>), ('dest', <tf.Tensor: shape=(2,), dtype=string,
numpy=array([b'SFO', b'SFO'], dtype=object)>), ('dep_hour', <tf.Tensor:
shape=(2,), dtype=int32, numpy=array([15, 1], dtype=int32)>), ('is_weekday',
<tf.Tensor: shape=(2,), dtype=int32, numpy=array([1, 1], dtype=int32)>),
('carrier', <tf.Tensor: shape=(2,), dtype=string, numpy=array([b'AA', b'OO'],
dtype=object)>), ('dep_airport_lat', <tf.Tensor: shape=(2,), dtype=float32,
numpy=array([41.979443, 34.42611 ], dtype=float32)>), ('dep_airport_lon',
<tf.Tensor: shape=(2,), dtype=float32, numpy=array([ -87.9075 , -119.84139],
dtype=float32)>), ('arr_airport_lat', <tf.Tensor: shape=(2,), dtype=float32,
numpy=array([37.61889, 37.61889], dtype=float32)>), ('arr_airport_lon',
<tf.Tensor: shape=(2,), dtype=float32, numpy=array([-122.375, -122.375],
dtype=float32)>)]), <tf.Tensor: shape=(2,), dtype=int32, numpy=array([0, 1],
dtype=int32)>)]

```

Create, train and evaluate TensorFlow model

```

[12]: real = {
    colname : tf.feature_column.numeric_column(colname)
    for colname in
        (
            'dep_delay,taxi_out,distance,dep_hour,is_weekday,' +
            'dep_airport_lat,dep_airport_lon,' +
            'arr_airport_lat,arr_airport_lon'
        ).split(',')
}
sparse = {

```

```

        'carrier': tf.feature_column.
        ↪categorical_column_with_vocabulary_list('carrier',
            vocabulary_list='AS,VX,F9,UA,US,WN,HA,EV,MQ,DL,OO,B6,NK,AA'.
        ↪split(',')),
        'origin' : tf.feature_column.
        ↪categorical_column_with_hash_bucket('origin', hash_bucket_size=1000),
        'dest'    : tf.feature_column.categorical_column_with_hash_bucket('dest', ↪
        ↪hash_bucket_size=1000),
    }

```

```

[13]: inputs = {
        colname : tf.keras.layers.Input(name=colname, shape=(), dtype='float32')
        for colname in real.keys()
    }
    inputs.update({
        colname : tf.keras.layers.Input(name=colname, shape=(), dtype='string')
        for colname in sparse.keys()
    })

```

### 0.3 Bucketing

```

[14]: latbuckets = np.linspace(20.0, 50.0, NBUCKETS).tolist() # USA
    lonbuckets = np.linspace(-120.0, -70.0, NBUCKETS).tolist() # USA
    disc = {}
    disc.update({
        'd_{}'.format(key) : tf.feature_column.bucketized_column(real[key], ↪
        ↪latbuckets)
        for key in ['dep_airport_lat', 'arr_airport_lat']
    })
    disc.update({
        'd_{}'.format(key) : tf.feature_column.bucketized_column(real[key], ↪
        ↪lonbuckets)
        for key in ['dep_airport_lon', 'arr_airport_lon']
    })
    # cross columns that make sense in combination
    sparse['dep_loc'] = tf.feature_column.crossed_column(
        [disc['d_dep_airport_lat'], disc['d_dep_airport_lon']], NBUCKETS*NBUCKETS)
    sparse['arr_loc'] = tf.feature_column.crossed_column(
        [disc['d_arr_airport_lat'], disc['d_arr_airport_lon']], NBUCKETS*NBUCKETS)
    sparse['dep_arr'] = tf.feature_column.crossed_column([sparse['dep_loc'], ↪
        ↪sparse['arr_loc']], NBUCKETS ** 4)
    # embed all the sparse columns
    embed = {
        'embed_{}'.format(colname) : tf.feature_column.embedding_column(col, ↪
        ↪NEMBEDS)
        for colname, col in sparse.items()
    }

```

```

}
real.update(embed)
# one-hot encode the sparse columns
sparse = {
    colname : tf.feature_column.indicator_column(col)
    for colname, col in sparse.items()
}
if DEVELOP_MODE:
    print(sparse.keys())
    print(real.keys())

```

```

dict_keys(['carrier', 'origin', 'dest', 'dep_loc', 'arr_loc', 'dep_arr'])
dict_keys(['dep_delay', 'taxi_out', 'distance', 'dep_hour', 'is_weekday',
'dep_airport_lat', 'dep_airport_lon', 'arr_airport_lat', 'arr_airport_lon',
'embed_carrier', 'embed_origin', 'embed_dest', 'embed_dep_loc', 'embed_arr_loc',
'embed_dep_arr'])

```

## 0.4 Train and evaluate the model

```

[15]: # Save the checkpoint:
output_dir='gs://{}/ch9/trained_model'.format(BUCKET)
os.environ['OUTDIR'] = output_dir # needed for deployment
print('Writing trained model to {}'.format(output_dir))

```

Writing trained model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model

```

[16]: # Delete the model checkpoints already present in the storage bucket:
!gsutil -m rm -rf $OUTDIR

```

CommandException: 1 files/objects could not be removed.

```

[17]: # Build a wide-and-deep model.
def wide_and_deep_classifier(inputs, linear_feature_columns,
    ↪dnn_feature_columns, dnn_hidden_units):
    deep = tf.keras.layers.DenseFeatures(dnn_feature_columns,
    ↪name='deep_inputs')(inputs)
    layers = [int(x) for x in dnn_hidden_units.split(',')]
    for layerno, numnodes in enumerate(layers):
        deep = tf.keras.layers.Dense(numnodes, activation='relu', name='dnn_{}'.
    ↪format(layerno+1))(deep)
    wide = tf.keras.layers.DenseFeatures(linear_feature_columns,
    ↪name='wide_inputs')(inputs)
    both = tf.keras.layers.concatenate([deep, wide], name='both')
    output = tf.keras.layers.Dense(1, activation='sigmoid', name='pred')(both)
    model = tf.keras.Model(inputs, output)
    model.compile(optimizer='adam',

```



```

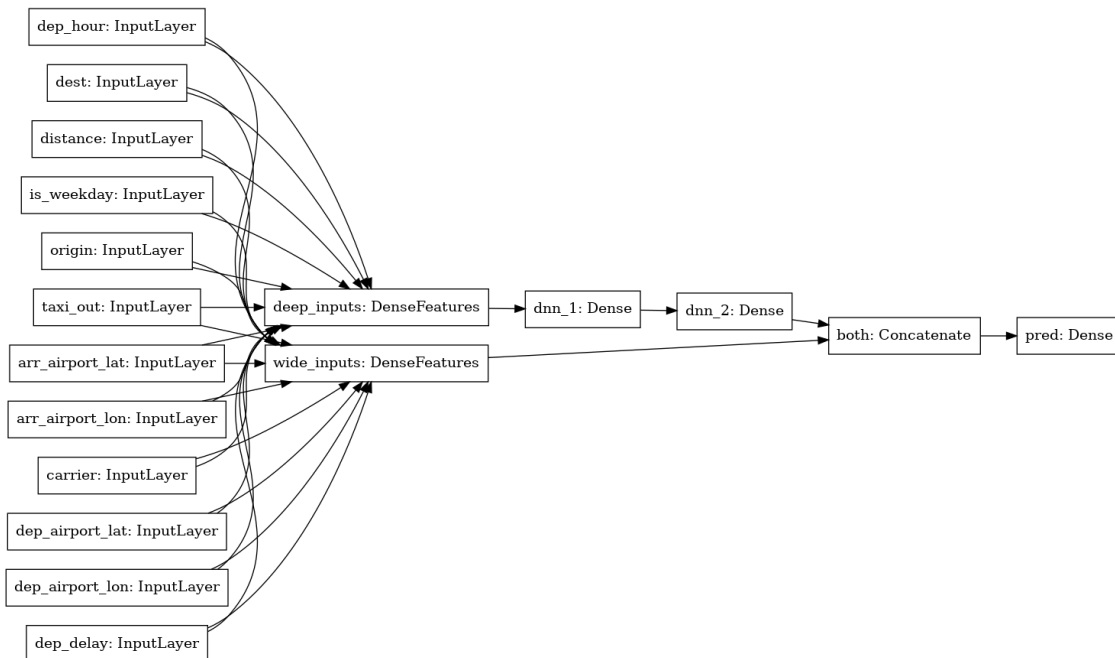
        loss='binary_crossentropy',
        metrics=['accuracy'])

    return model

model = wide_and_deep_classifier(
    inputs,
    linear_feature_columns = sparse.values(),
    dnn_feature_columns = real.values(),
    dnn_hidden_units = DNN_HIDDEN_UNITS)
tf.keras.utils.plot_model(model, 'flights_model.png', show_shapes=False,
    rankdir='LR')

```

[17]:



```

[18]: # training and evaluation dataset
train_batch_size = TRAIN_BATCH_SIZE
if DEVELOP_MODE:
    eval_batch_size = 100
    steps_per_epoch = 3
    epochs = 50
    num_eval_examples = eval_batch_size*10
else:
    eval_batch_size = 100
    steps_per_epoch = NUM_EXAMPLES // train_batch_size
    epochs = 10
    num_eval_examples = eval_batch_size * 100
train_dataset = read_dataset(training_data_uri, train_batch_size)

```

```

eval_dataset = read_dataset(validation_data_uri, eval_batch_size, tf.estimator.
    ↪ModeKeys.EVAL, num_eval_examples)
checkpoint_path = '{}/checkpoints/flights.cpt'.format(output_dir)
shutil.rmtree(checkpoint_path, ignore_errors=True)
cp_callback = tf.keras.callbacks.ModelCheckpoint(checkpoint_path,
                                                save_weights_only=True,
                                                verbose=1)

history = model.fit(train_dataset,
                    validation_data=eval_dataset,
                    epochs=epochs,
                    steps_per_epoch=steps_per_epoch,
                    callbacks=[cp_callback])

```

Epoch 1/50

3/3 [=====] - 10s 3s/step - loss: 1.8575 - accuracy: 0.7656 - val\_loss: 0.7295 - val\_accuracy: 0.6385

Epoch 00001: saving model to gs://qwiklabs-

gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt

Epoch 2/50

3/3 [=====] - 5s 3s/step - loss: 0.7876 - accuracy: 0.4844 - val\_loss: 0.6990 - val\_accuracy: 0.5778

Epoch 00002: saving model to gs://qwiklabs-

gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt

Epoch 3/50

3/3 [=====] - 5s 2s/step - loss: 0.7202 - accuracy: 0.7240 - val\_loss: 0.5791 - val\_accuracy: 0.7945

Epoch 00003: saving model to gs://qwiklabs-

gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt

Epoch 4/50

3/3 [=====] - 5s 2s/step - loss: 0.5792 - accuracy: 0.8021 - val\_loss: 0.4814 - val\_accuracy: 0.8128

Epoch 00004: saving model to gs://qwiklabs-

gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt

Epoch 5/50

3/3 [=====] - 4s 2s/step - loss: 0.4880 - accuracy: 0.8125 - val\_loss: 0.4504 - val\_accuracy: 0.8547

Epoch 00005: saving model to gs://qwiklabs-

gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt

Epoch 6/50

3/3 [=====] - 10s 5s/step - loss: 0.4705 - accuracy: 0.7917 - val\_loss: 0.3994 - val\_accuracy: 0.8598

Epoch 00006: saving model to gs://qwiklabs-gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 7/50  
3/3 [=====] - 5s 2s/step - loss: 0.3749 - accuracy: 0.8854 - val\_loss: 0.3681 - val\_accuracy: 0.8582

Epoch 00007: saving model to gs://qwiklabs-gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 8/50  
3/3 [=====] - 5s 2s/step - loss: 0.3520 - accuracy: 0.8542 - val\_loss: 0.3544 - val\_accuracy: 0.8698

Epoch 00008: saving model to gs://qwiklabs-gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 9/50  
3/3 [=====] - 5s 2s/step - loss: 0.3478 - accuracy: 0.8646 - val\_loss: 0.3459 - val\_accuracy: 0.8763

Epoch 00009: saving model to gs://qwiklabs-gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 10/50  
3/3 [=====] - 5s 3s/step - loss: 0.4070 - accuracy: 0.8229 - val\_loss: 0.3329 - val\_accuracy: 0.8755

Epoch 00010: saving model to gs://qwiklabs-gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 11/50  
3/3 [=====] - 5s 2s/step - loss: 0.3577 - accuracy: 0.8490 - val\_loss: 0.3635 - val\_accuracy: 0.8582

Epoch 00011: saving model to gs://qwiklabs-gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 12/50  
3/3 [=====] - 5s 2s/step - loss: 0.3561 - accuracy: 0.8750 - val\_loss: 0.3186 - val\_accuracy: 0.8702

Epoch 00012: saving model to gs://qwiklabs-gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 13/50  
3/3 [=====] - 5s 2s/step - loss: 0.3860 - accuracy: 0.8125 - val\_loss: 0.3099 - val\_accuracy: 0.8751

Epoch 00013: saving model to gs://qwiklabs-gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 14/50  
3/3 [=====] - 5s 2s/step - loss: 0.3037 - accuracy: 0.8646 - val\_loss: 0.3639 - val\_accuracy: 0.8570

Epoch 00014: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 15/50  
3/3 [=====] - 5s 2s/step - loss: 0.3168 - accuracy:  
0.8854 - val\_loss: 0.2979 - val\_accuracy: 0.8973

Epoch 00015: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 16/50  
3/3 [=====] - 5s 3s/step - loss: 0.2573 - accuracy:  
0.8958 - val\_loss: 0.3173 - val\_accuracy: 0.8734

Epoch 00016: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 17/50  
3/3 [=====] - 4s 2s/step - loss: 0.2906 - accuracy:  
0.8854 - val\_loss: 0.2796 - val\_accuracy: 0.8964

Epoch 00017: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 18/50  
3/3 [=====] - 5s 2s/step - loss: 0.2949 - accuracy:  
0.8906 - val\_loss: 0.3319 - val\_accuracy: 0.8863

Epoch 00018: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 19/50  
3/3 [=====] - 5s 2s/step - loss: 0.2946 - accuracy:  
0.9010 - val\_loss: 0.2827 - val\_accuracy: 0.8897

Epoch 00019: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 20/50  
3/3 [=====] - 4s 2s/step - loss: 0.2519 - accuracy:  
0.8958 - val\_loss: 0.3133 - val\_accuracy: 0.8816

Epoch 00020: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 21/50  
3/3 [=====] - 5s 3s/step - loss: 0.2486 - accuracy:  
0.8906 - val\_loss: 0.2614 - val\_accuracy: 0.9092

Epoch 00021: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 22/50  
3/3 [=====] - 5s 2s/step - loss: 0.2917 - accuracy:  
0.8958 - val\_loss: 0.3318 - val\_accuracy: 0.8902

Epoch 00022: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 23/50  
3/3 [=====] - 5s 2s/step - loss: 0.3465 - accuracy:  
0.8906 - val\_loss: 0.3170 - val\_accuracy: 0.8853

Epoch 00023: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 24/50  
3/3 [=====] - 5s 2s/step - loss: 0.3833 - accuracy:  
0.8802 - val\_loss: 0.2575 - val\_accuracy: 0.9057

Epoch 00024: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 25/50  
3/3 [=====] - 5s 3s/step - loss: 0.2323 - accuracy:  
0.9323 - val\_loss: 0.3078 - val\_accuracy: 0.9161

Epoch 00025: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 26/50  
3/3 [=====] - 5s 3s/step - loss: 0.2925 - accuracy:  
0.9115 - val\_loss: 0.2711 - val\_accuracy: 0.8996

Epoch 00026: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 27/50  
3/3 [=====] - 5s 2s/step - loss: 0.3276 - accuracy:  
0.8906 - val\_loss: 0.2900 - val\_accuracy: 0.8942

Epoch 00027: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 28/50  
3/3 [=====] - 5s 2s/step - loss: 0.2107 - accuracy:  
0.9219 - val\_loss: 0.2527 - val\_accuracy: 0.9263

Epoch 00028: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 29/50  
3/3 [=====] - 5s 3s/step - loss: 0.2910 - accuracy:  
0.9010 - val\_loss: 0.2892 - val\_accuracy: 0.9154

Epoch 00029: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 30/50  
3/3 [=====] - 5s 3s/step - loss: 0.2698 - accuracy:  
0.9167 - val\_loss: 0.2330 - val\_accuracy: 0.9172

Epoch 00030: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 31/50  
3/3 [=====] - 5s 2s/step - loss: 0.1780 - accuracy:  
0.9427 - val\_loss: 0.2685 - val\_accuracy: 0.9011

Epoch 00031: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 32/50  
3/3 [=====] - 5s 2s/step - loss: 0.2694 - accuracy:  
0.8958 - val\_loss: 0.2229 - val\_accuracy: 0.9237

Epoch 00032: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 33/50  
3/3 [=====] - 10s 5s/step - loss: 0.2536 - accuracy:  
0.9115 - val\_loss: 0.2601 - val\_accuracy: 0.9254

Epoch 00033: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 34/50  
3/3 [=====] - 5s 3s/step - loss: 0.2541 - accuracy:  
0.9375 - val\_loss: 0.2240 - val\_accuracy: 0.9196

Epoch 00034: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 35/50  
3/3 [=====] - 5s 2s/step - loss: 0.1745 - accuracy:  
0.9271 - val\_loss: 0.2295 - val\_accuracy: 0.9148

Epoch 00035: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 36/50  
3/3 [=====] - 5s 3s/step - loss: 0.2303 - accuracy:  
0.9219 - val\_loss: 0.2293 - val\_accuracy: 0.9281

Epoch 00036: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 37/50  
3/3 [=====] - 4s 2s/step - loss: 0.2425 - accuracy:  
0.9062 - val\_loss: 0.2207 - val\_accuracy: 0.9282

Epoch 00037: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 38/50  
3/3 [=====] - 4s 2s/step - loss: 0.2131 - accuracy:  
0.9219 - val\_loss: 0.2264 - val\_accuracy: 0.9165

Epoch 00038: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 39/50  
3/3 [=====] - 5s 2s/step - loss: 0.2179 - accuracy:  
0.9167 - val\_loss: 0.2086 - val\_accuracy: 0.9267

Epoch 00039: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 40/50  
3/3 [=====] - 5s 2s/step - loss: 0.2123 - accuracy:  
0.9271 - val\_loss: 0.2104 - val\_accuracy: 0.9326

Epoch 00040: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 41/50  
3/3 [=====] - 5s 2s/step - loss: 0.1402 - accuracy:  
0.9740 - val\_loss: 0.2220 - val\_accuracy: 0.9217

Epoch 00041: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 42/50  
3/3 [=====] - 5s 2s/step - loss: 0.2652 - accuracy:  
0.9010 - val\_loss: 0.2034 - val\_accuracy: 0.9340

Epoch 00042: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 43/50  
3/3 [=====] - 5s 2s/step - loss: 0.3511 - accuracy:  
0.8385 - val\_loss: 0.2029 - val\_accuracy: 0.9343

Epoch 00043: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 44/50  
3/3 [=====] - 4s 2s/step - loss: 0.2008 - accuracy:  
0.9115 - val\_loss: 0.3299 - val\_accuracy: 0.9034

Epoch 00044: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 45/50  
3/3 [=====] - 5s 2s/step - loss: 0.3458 - accuracy:  
0.8906 - val\_loss: 0.1948 - val\_accuracy: 0.9309

Epoch 00045: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/checkpoints/flights.cpt  
Epoch 46/50  
3/3 [=====] - 4s 2s/step - loss: 0.2466 - accuracy:  
0.8958 - val\_loss: 0.2215 - val\_accuracy: 0.9250

```
Epoch 00046: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained_model/checkpoints/flights.cpt  
Epoch 47/50  
3/3 [=====] - 4s 2s/step - loss: 0.2470 - accuracy:  
0.9219 - val_loss: 0.2642 - val_accuracy: 0.9140
```

```
Epoch 00047: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained_model/checkpoints/flights.cpt  
Epoch 48/50  
3/3 [=====] - 5s 2s/step - loss: 0.2612 - accuracy:  
0.9062 - val_loss: 0.1900 - val_accuracy: 0.9341
```


```
Epoch 00048: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained_model/checkpoints/flights.cpt  
Epoch 49/50  
3/3 [=====] - 5s 2s/step - loss: 0.1945 - accuracy:  
0.9219 - val_loss: 0.2393 - val_accuracy: 0.9209
```

```
Epoch 00049: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained_model/checkpoints/flights.cpt  
Epoch 50/50  
3/3 [=====] - 5s 2s/step - loss: 0.1732 - accuracy:  
0.9427 - val_loss: 0.1918 - val_accuracy: 0.9350
```

```
Epoch 00050: saving model to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained_model/checkpoints/flights.cpt
```

```
[19]: import matplotlib.pyplot as plt  
nrows = 1  
ncols = 2  
fig = plt.figure(figsize=(10, 5))  
for idx, key in enumerate(['loss', 'accuracy']):  
    ax = fig.add_subplot(nrows, ncols, idx+1)  
    plt.plot(history.history[key])  
    plt.plot(history.history['val_{}'.format(key)])  
    plt.title('model {}'.format(key))  
    plt.ylabel(key)  
    plt.xlabel('epoch')  
    plt.legend(['train', 'validation'], loc='upper left');
```





## 0.5 Export the trained model

```
[20]: import time
export_dir = '{}'/export/flights_{}'.format(output_dir, time.
    strftime("%Y%m%d-%H%M%S"))
print('Exporting to {}'.format(export_dir))
tf.saved_model.save(model, export_dir)
```

Exporting to gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/export/flights\_20230425-184304

2023-04-25 18:43:07.715159: W tensorflow/python/util/util.cc:348] Sets are not currently considered sequences, but this may change in the future, so consider avoiding using them.

INFO:tensorflow:Assets written to: gs://qwiklabs-gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/export/flights\_20230425-184304/assets

## 0.6 Deploy flights model to Vertex AI

```
[21]: %bash
# note TF_VERSION and ENDPOINT_NAME set in 1st cell
# TF_VERSION=2-6
# ENDPOINT_NAME=flights
TIMESTAMP=$(date +%Y%m%d-%H%M%S)
MODEL_NAME=${ENDPOINT_NAME}-${TIMESTAMP}
EXPORT_PATH=$(gsutil ls ${OUTDIR}/export | tail -1)
echo $EXPORT_PATH
```

```

# create the model endpoint for deploying the model
if [[ $(gcloud beta ai endpoints list --region=$REGION \
    --format='value(DISPLAY_NAME)' --filter=display_name=${ENDPOINT_NAME})
↪]]; then
    echo "Endpoint for $MODEL_NAME already exists"
else
    echo "Creating Endpoint for $MODEL_NAME"
    gcloud beta ai endpoints create --region=${REGION}
↪--display-name=${ENDPOINT_NAME}
fi
ENDPOINT_ID=$(gcloud beta ai endpoints list --region=$REGION \
    --format='value(ENDPOINT_ID)'
↪--filter=display_name=${ENDPOINT_NAME})
echo "ENDPOINT_ID=$ENDPOINT_ID"
# delete any existing models with this name
for MODEL_ID in $(gcloud beta ai models list --region=$REGION
↪--format='value(MODEL_ID)' --filter=display_name=${MODEL_NAME}); do
    echo "Deleting existing $MODEL_NAME ... $MODEL_ID "
    gcloud ai models delete --region=$REGION $MODEL_ID
done
# create the model using the parameters docker container image and artifact uri
gcloud beta ai models upload --region=$REGION --display-name=$MODEL_NAME \
    --container-image-uri=us-docker.pkg.dev/vertex-ai/prediction/tf2-cpu.
↪${TF_VERSION}:latest \
    --artifact-uri=$EXPORT_PATH
MODEL_ID=$(gcloud beta ai models list --region=$REGION
↪--format='value(MODEL_ID)' --filter=display_name=${MODEL_NAME})
echo "MODEL_ID=$MODEL_ID"
# deploy the model to the endpoint
gcloud beta ai endpoints deploy-model $ENDPOINT_ID \
    --region=$REGION \
    --model=$MODEL_ID \
    --display-name=$MODEL_NAME \
    --machine-type=n1-standard-2 \
    --min-replica-count=1 \
    --max-replica-count=1 \
    --traffic-split=0=100

```

gs://qwiklabs-

gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/export/flights\_20230425-184304/

Creating Endpoint for flights-20230425-184315

ENDPOINT\_ID=366388060700540928

MODEL\_ID=5491376684508643328

Using endpoint [https://us-central1-aiplatform.googleapis.com/]

WARNING: The following filter keys were not present in any resource :  
display\_name

Using endpoint [https://us-central1-aiplatform.googleapis.com/]  
Waiting for operation [1087355221060878336]...

..done.

Using endpoint [https://us-central1-aiplatform.googleapis.com/]

WARNING: The following filter keys were not present in any resource :  
display\_name

Waiting for operation [7298945007111634944]...

```
..done.
```

Using endpoint [https://us-central1-aiplatform.googleapis.com/]

• • •  
• • •  
• • •  
• • •  
• • •  
• • •  
• • •

```
..done.
```

Deployed a model to the endpoint 366388060700540928. Id of the deployed model: 3791399766571614208.

Create a test input file

```
[22]: %%writefile example_input.json
      {"instances": [
```

```
{
  "dep_hour": 2, "is_weekday": 1, "dep_delay": 40, "taxi_out": 17, "distance": 41,
  "carrier": "AS", "dep_airport_lat": 58.42527778, "dep_airport_lon": -135.7075,
  "arr_airport_lat": 58.35472222, "arr_airport_lon": -134.57472222,
  "origin": "GST", "dest": "JNU"},
  {
    "dep_hour": 22, "is_weekday": 0, "dep_delay": -7, "taxi_out": 7, "distance": 201,
    "carrier": "HA", "dep_airport_lat": 21.97611111, "dep_airport_lon": -159.33888889,
    "arr_airport_lat": 20.89861111, "arr_airport_lon": -156.43055556,
    "origin": "LIH", "dest": "OGG"}
}
```

Writing example\_input.json

Make a prediction from the model endpoint.

```
[23]: % bash
ENDPOINT_ID=$(gcloud beta ai endpoints list --region=$REGION \
        --format='value(ENDPOINT_ID)' \
        --filter=display_name=${ENDPOINT_NAME})
echo $ENDPOINT_ID
gcloud beta ai endpoints predict $ENDPOINT_ID --region=$REGION \
        --json-request=example_input.json
```

```
366388060700540928
```

```
[[0.0247025788], [0.983607709]]
```

Using endpoint [https://us-central1-aiplatform.googleapis.com/]

Using endpoint [https://us-central1-prediction-aiplatform.googleapis.com/]

Send an HTTP POST request and you will get the result back as JSON:

```
[24]: % bash
PROJECT=$(gcloud config get-value project)
ENDPOINT_ID=$(gcloud beta ai endpoints list --region=$REGION \
        --format='value(ENDPOINT_ID)' \
        --filter=display_name=${ENDPOINT_NAME})
curl -X POST \
  -H "Authorization: Bearer $(gcloud auth application-default \
    print-access-token) \
  -H "Content-Type: application/json; charset=utf-8" \
  -d @example_input.json \
  "https://${REGION}-aiplatform.googleapis.com/v1/projects/${PROJECT}/locations/
  ${REGION}/endpoints/${ENDPOINT_ID}:predict"
```

```
{
  "predictions": [
    [
      0.0247025788
    ],
    [

```

```

    0.983607709
  ]
],
"deployedModelId": "3791399766571614208",
"model": "projects/691390126128/locations/us-
central1/models/5491376684508643328",
"modelDisplayName": "flights-20230425-184315",
"modelVersionId": "1"
}

```

Using endpoint [https://us-central1-aiplatform.googleapis.com/]

% Total	% Received	% Xferd	Average Speed	Time	Time	Time	Current
			Dload Upload	Total	Spent	Left	Speed
100	835	0 292 100	543 4709 8758	--:--:--	--:--:--	--:--:--	13467

## 0.7 Model explainability

```

[25]: %%%bash
model_dir=$(gsutil ls ${OUTDIR}/export | tail -1)
echo $model_dir
saved_model_cli show --tag_set serve --signature_def serving_default --dir_
↪$model_dir

```

gs://qwiklabs-  
gcp-01-c3764cc81f6b-dsongcp/ch9/trained\_model/export/flights\_20230425-184304/

The given SavedModel SignatureDef contains the following input(s):

```

inputs['arr_airport_lat'] tensor_info:
  dtype: DT_FLOAT
  shape: (-1)
  name: serving_default_arr_airport_lat:0
inputs['arr_airport_lon'] tensor_info:
  dtype: DT_FLOAT
  shape: (-1)
  name: serving_default_arr_airport_lon:0
inputs['carrier'] tensor_info:
  dtype: DT_STRING
  shape: (-1)
  name: serving_default_carrier:0
inputs['dep_airport_lat'] tensor_info:
  dtype: DT_FLOAT
  shape: (-1)
  name: serving_default_dep_airport_lat:0
inputs['dep_airport_lon'] tensor_info:
  dtype: DT_FLOAT
  shape: (-1)
  name: serving_default_dep_airport_lon:0
inputs['dep_delay'] tensor_info:
  dtype: DT_FLOAT

```

```

        shape: (-1)
        name: serving_default_dep_delay:0
inputs['dep_hour'] tensor_info:
    dtype: DT_FLOAT
    shape: (-1)
    name: serving_default_dep_hour:0
inputs['dest'] tensor_info:
    dtype: DT_STRING
    shape: (-1)
    name: serving_default_dest:0
inputs['distance'] tensor_info:
    dtype: DT_FLOAT
    shape: (-1)
    name: serving_default_distance:0
inputs['is_weekday'] tensor_info:
    dtype: DT_FLOAT
    shape: (-1)
    name: serving_default_is_weekday:0
inputs['origin'] tensor_info:
    dtype: DT_STRING
    shape: (-1)
    name: serving_default_origin:0
inputs['taxi_out'] tensor_info:
    dtype: DT_FLOAT
    shape: (-1)
    name: serving_default_taxi_out:0
The given SavedModel SignatureDef contains the following output(s):
outputs['pred'] tensor_info:
    dtype: DT_FLOAT
    shape: (-1, 1)
    name: StatefulPartitionedCall_2:0
Method name is: tensorflow/serving/predict

```

```

[26]: cols = ('dep_delay,taxi_out,distance,dep_hour,is_weekday,' +
             'dep_airport_lat,dep_airport_lon,' +
             'arr_airport_lat,arr_airport_lon,' +
             'carrier,origin,dest')
inputs = {x: {"inputTensorName": "{}".format(x)}
          for x in cols.split(',') }
expl = {
    "inputs": inputs,
    "outputs": {
        "pred": {
            "outputTensorName": "pred"
        }
    }
}

```

```
print(expl)
with open('explanation-metadata.json', 'w') as ofp:
    json.dump(expl, ofp, indent=2)
```

```
{'inputs': {'dep_delay': {'inputTensorName': 'dep_delay'}, 'taxi_out':
{'inputTensorName': 'taxi_out'}, 'distance': {'inputTensorName': 'distance'},
'dep_hour': {'inputTensorName': 'dep_hour'}, 'is_weekday': {'inputTensorName':
'is_weekday'}, 'dep_airport_lat': {'inputTensorName': 'dep_airport_lat'},
'dep_airport_lon': {'inputTensorName': 'dep_airport_lon'}, 'arr_airport_lat':
{'inputTensorName': 'arr_airport_lat'}, 'arr_airport_lon': {'inputTensorName':
'arr_airport_lon'}, 'carrier': {'inputTensorName': 'carrier'}, 'origin':
{'inputTensorName': 'origin'}, 'dest': {'inputTensorName': 'dest'}}, 'outputs':
{'pred': {'outputTensorName': 'pred'}}}
```

```
[27]: !cat explanation-metadata.json
```

```
{
  "inputs": {
    "dep_delay": {
      "inputTensorName": "dep_delay"
    },
    "taxi_out": {
      "inputTensorName": "taxi_out"
    },
    "distance": {
      "inputTensorName": "distance"
    },
    "dep_hour": {
      "inputTensorName": "dep_hour"
    },
    "is_weekday": {
      "inputTensorName": "is_weekday"
    },
    "dep_airport_lat": {
      "inputTensorName": "dep_airport_lat"
    },
    "dep_airport_lon": {
      "inputTensorName": "dep_airport_lon"
    },
    "arr_airport_lat": {
      "inputTensorName": "arr_airport_lat"
    },
    "arr_airport_lon": {
      "inputTensorName": "arr_airport_lon"
    },
    "carrier": {
      "inputTensorName": "carrier"
    },
  },
```



```

    "origin": {
      "inputTensorName": "origin"
    },
    "dest": {
      "inputTensorName": "dest"
    }
  },
  "outputs": {
    "pred": {
      "outputTensorName": "pred"
    }
  }
}
}

```

## 0.8 Create and deploy another model flights\_xai to Vertex AI

```

[28]: %%%bash
# note TF_VERSION set in 1st cell, but ENDPOINT_NAME is being changed
# TF_VERSION=2-6
ENDPOINT_NAME=flights_xai
TIMESTAMP=$(date +%Y%m%d-%H%M%S)
MODEL_NAME=${ENDPOINT_NAME}-${TIMESTAMP}
EXPORT_PATH=$(gsutil ls ${OUTDIR}/export | tail -1)
echo $EXPORT_PATH
# create the model endpoint for deploying the model
if [[ $(gcloud beta ai endpoints list --region=$REGION \
      --format='value(DISPLAY_NAME)' --filter=display_name=${ENDPOINT_NAME})
↪]]; then
  echo "Endpoint for $MODEL_NAME already exists"
else
  # create model endpoint
  echo "Creating Endpoint for $MODEL_NAME"
  gcloud beta ai endpoints create --region=${REGION}
↪--display-name=${ENDPOINT_NAME}
fi
ENDPOINT_ID=$(gcloud beta ai endpoints list --region=$REGION \
      --format='value(ENDPOINT_ID)'
↪--filter=display_name=${ENDPOINT_NAME})
echo "ENDPOINT_ID=$ENDPOINT_ID"
# delete any existing models with this name
for MODEL_ID in $(gcloud beta ai models list --region=$REGION
↪--format='value(MODEL_ID)' --filter=display_name=${MODEL_NAME}); do
  echo "Deleting existing $MODEL_NAME ... $MODEL_ID "
  gcloud ai models delete --region=$REGION $MODEL_ID
done
# upload the model using the parameters docker container image, artifact URI,
↪explanation method,

```

```
# explanation path count and explanation metadata JSON file
↪`explanation-metadata.json`.
# Here, you keep number of feature permutations to `10` when approximating the
↪Shapley values for explanation.
gcloud beta ai models upload --region=$REGION --display-name=$MODEL_NAME \
  --container-image-uri=us-docker.pkg.dev/vertex-ai/prediction/tf2-cpu.
↪${TF_VERSION}:latest \
  --artifact-uri=$EXPORT_PATH \
  --explanation-method=sampled-shapley --explanation-path-count=10
↪--explanation-metadata-file=explanation-metadata.json
MODEL_ID=$(gcloud beta ai models list --region=$REGION
↪--format='value(MODEL_ID)' --filter=display_name=${MODEL_NAME})
echo "MODEL_ID=$MODEL_ID"
# deploy the model to the endpoint
gcloud beta ai endpoints deploy-model $ENDPOINT_ID \
  --region=$REGION \
  --model=$MODEL_ID \
  --display-name=$MODEL_NAME \
  --machine-type=n1-standard-2 \
  --min-replica-count=1 \
  --max-replica-count=1 \
  --traffic-split=0=100
```

```
gs://qwiklabs-
gcp-01-c3764cc81f6b-dsongcp/ch9/trained_model/export/flights_20230425-184304/
Creating Endpoint for flights_xai-20230425-190300
ENDPOINT_ID=419305356322144256
MODEL_ID=3185533675294949376

Using endpoint [https://us-central1-aiplatform.googleapis.com/]
Using endpoint [https://us-central1-aiplatform.googleapis.com/]
Waiting for operation [8182776433983094784]...
...done.
Created Vertex AI endpoint: projects/691390126128/locations/us-
central1/endpoints/419305356322144256.
Using endpoint [https://us-central1-aiplatform.googleapis.com/]
Using endpoint [https://us-central1-aiplatform.googleapis.com/]
Using endpoint [https://us-central1-aiplatform.googleapis.com/]
Waiting for operation [237300791394697216]...
...
...
...
...
...
...
...
...
```

...  
...  
...  
...  
...  
...  
...  
...  
...  
...  
...  
...  
...  
...C  
Us  
Us  
Wa  
...  
...  
...  
...  
...  
...  
...  
...  
...  
...  
...

```
..done.
```

Using endpoint [https://us-central1-aiplatform.googleapis.com/]

Using endpoint [https://us-central1-aiplatform.googleapis.com/]

```
Waiting for operation [3789514997483175936]...
```

...

...done.

Deployed a model to the endpoint 419305356322144256. Id of the deployed model: 4363356919247667200.

```
[29]: %%%bash
PROJECT=$(gcloud config get-value project)
ENDPOINT_NAME=flights_xai
ENDPOINT_ID=$(gcloud beta ai endpoints list --region=$REGION \
               --format='value(ENDPOINT_ID)' \
               --filter=display_name=${ENDPOINT_NAME})
curl -X POST \
  -H "Authorization: Bearer "$(gcloud auth application-default \
  --print-access-token) \
  -H "Content-Type: application/json; charset=utf-8" \
  -d @example_input.json \
  "https://$REGION-aiplatform.googleapis.com/v1/projects/${PROJECT}/locations/
  $REGION/endpoints/${ENDPOINT_ID}:explain"
```

```
{
  "error": {
    "code": 400,
    "message": "\"Explainability failed with exception: Exceeded max_retries (5)
while Explainer attempting to call Predictor. Error: \u003c_InactiveRpcError of
RPC that terminated with:\\n\\tstatus =
StatusCode.RESOURCE_EXHAUSTED\\n\\tdetails = \\\"\\\"\\n\\tdebug_error_string =
\\\"UNKNOWN:Error received from peer ipv4:10.60.9.2:8500
{created_time:\\\"2023-04-25T19:17:28.014529189+00:00\\\", grpc_status:8,
grpc_message:\\\"\\\"}\\\"\\\"\\n\\u003e\\\"\\n\",
    \"status\": \"FAILED_PRECONDITION\"
  }
}
```

Using endpoint [https://us-central1-aiplatform.googleapis.com/]

% Total	% Received	% Xferd	Average Speed	Time	Time	Time	Current
			Dload Upload	Total	Spent	Left	Speed
100	1071	0 528 100	543	213 219	0:00:02 0:00:02	--:--:--	432

[ ]: